

# Architectures

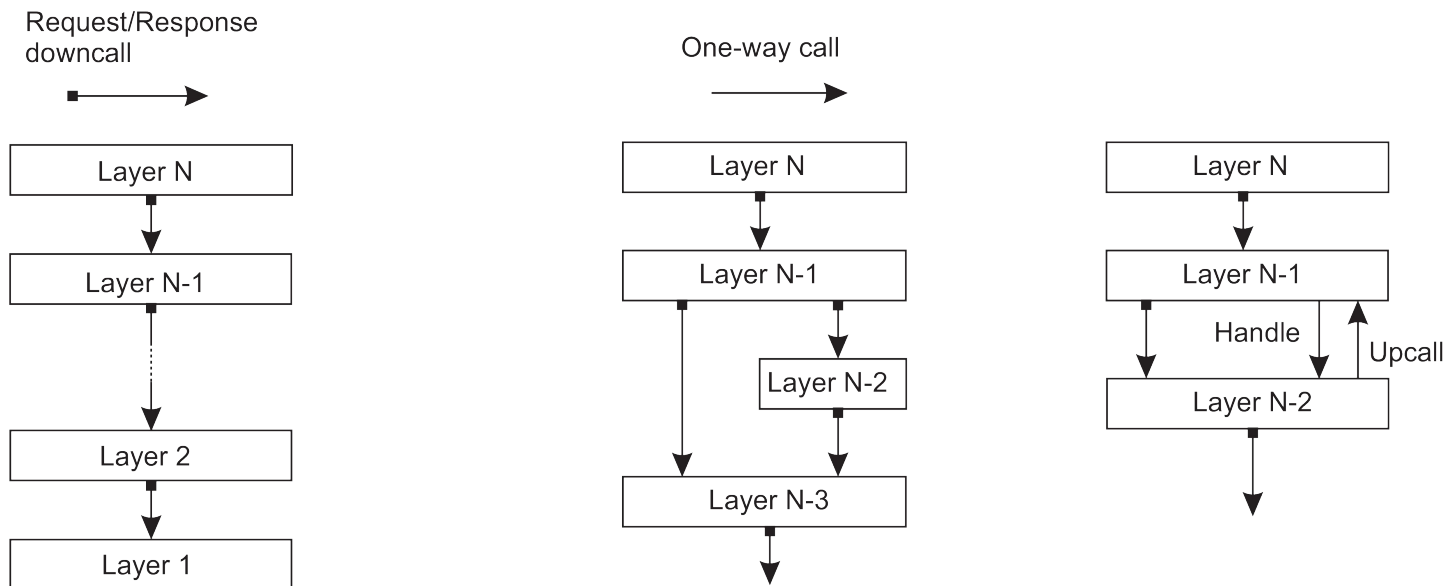
---

Dr. Padraig Corcoran

- A **software architecture** concerns the logical organisation of the software components.
- Important software architectures in distributed systems are:
  - Layered architecture.
  - Object-based architecture.
  - Service-Oriented Architecture (SOA).
- A **system architecture** concerns where these software components are placed across the various machines.
- Important system architectures in distributed systems are:
  - Centralised client-server architecture.
  - Decentralised peer-to-peer architecture.
  - Hybrid architecture.
  - Edge computing

# Layered architectures


- Components are organized in a layered fashion.
- A component at layer  $j$  can make a downcall to a component at a lower-level layer  $i$  (with  $i < j$ ).
- Only in exceptional cases will an upcall be made.



(a) *Pure layered organization.* (b) *Mixed layered organisation.*  
(c) *Layered organization with upcalls*

# Object-based architecture

- Objects provide a way of *encapsulating* data and operations that can be performed on that data.
- Each object corresponds to a software component and these are connected through method calls.




*An object-based architectural style.*

# Service-Oriented Architecture (SOA)

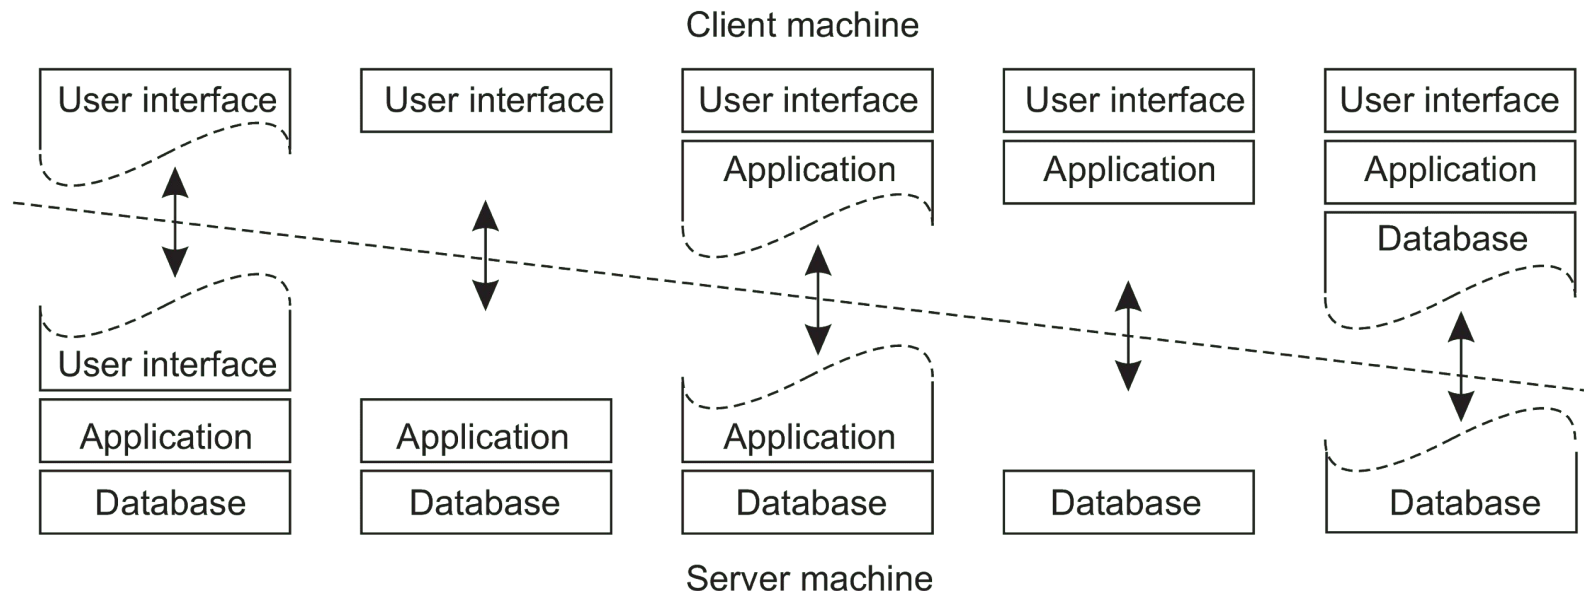
- A distributed system is constructed as a composition of different services (e.g. payment and database services).
- Developing a distributed system involves service composition or mashup.
- Each service must offer a well-defined (programming) interface.
- Can be implemented using Web services (e.g. SOAP and RESTful web services).

# Centralized client-server architecture

- Processes divided into two (possibly overlapping) groups.
- A client is a process that requests a service from a server by sending it a request and subsequently waiting for reply.
- This client-server interaction, also known as **request-reply behaviour**, is shown.




*General interaction between a client and a server.*



*Thin and fat clients are possible.*

- A server may sometimes need to act as a client.




*An example of a server acting as client.*




# Network topology


- A graph describing the layout of the distributed system.
- Nodes can correspond to physical devices or processes.
- Edges correspond to communication channels.



(a)




(b)




(c)

*Examples of network topologies: (a) ring, (b) directed tree and (c) 3D cube.*


- The topology is **structured** if it “has a pattern”.
- The topology is **unstructured** if the edges are random.
- In unstructured networks many tasks, such as routing, are more complicated.



(a)




(b)




(c)

*Structured topology*



*Unstructured topology*


- If the nodes correspond to processes, the graph is commonly referred to as an **overlay network**.
- Nodes in the overlay network can be thought of as being connected by virtual or logical links, each of which corresponds to a path in the underlying network.




*Example of overlay network.*

# Decentralized peer-to-peer (P2P) architecture

- **Vertical distribution** is characterised by having logically different processes (e.g. client-server model).
- **Horizontal distribution** is characterised by having logically equal processes (e.g. peer-to-peer model).
- The processes in a peer-to-peer system are all equal (will act as both client and server).
- The peer-to-peer architecture is commonly employed in file sharing systems.




*Client-server model*




*Peer-to-peer model*


- In a **structured peer-to-peer architecture** the network topology is structured.
- In an **unstructured peer-to-peer architecture** the network topology is unstructured.



(a)




(b)



(c)


*Structured topology*



*Unstructured topology*

# Locating data in a P2P system

- A given node wants to locate/access data stored on a different node.
- The node seeking the data only knows the **key** corresponding to the data in question.
- Problem - given a key locate the node containing the corresponding data.



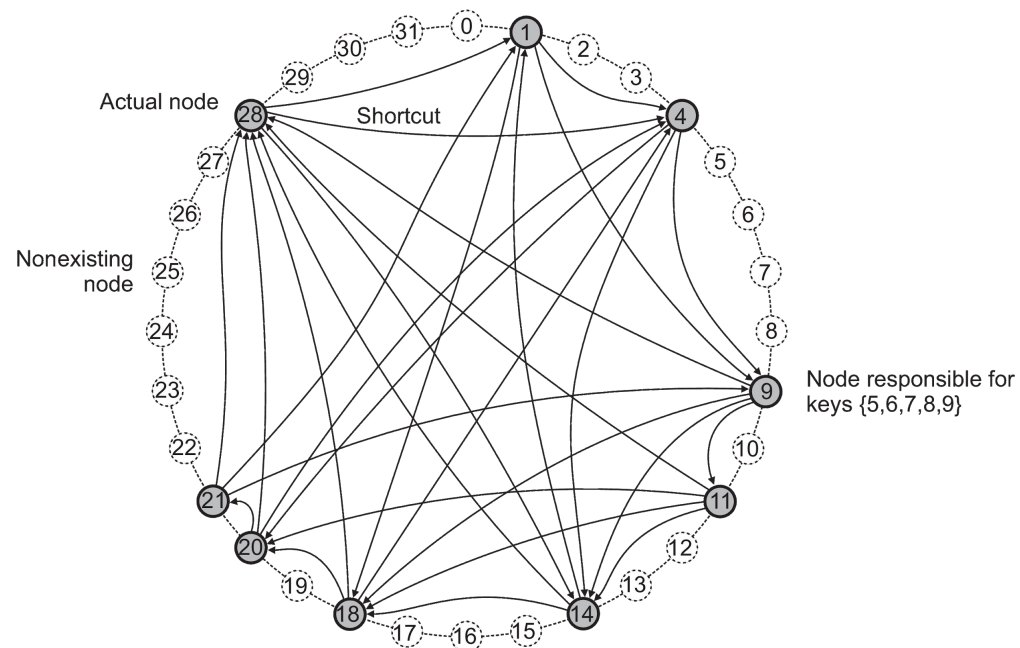
- Keys are often determined using a hash function.
- A hash function  $H(.)$  reduces a variable length input to a fixed length output.
- A hash function will have the following properties:
  - Computed  $m = H(M)$  is easy but computing the inverse is impossible.
  - If  $M \neq M'$  then  $H(M) \neq H(M')$ .




- In structured peer-to-peer systems, one can efficiently determine the location of data (mapping from **data key** to **data location**).
- Algorithms include **Chord**.
- In unstructured peer-to-peer system this is not the case and instead searching must be performed.
- Searching algorithms include **flooding** and **random walk**.

# Searching in structured P2P using Chord


- Nodes organized in a ring topology.
- Each node is assigned an *identifier* number.
- Each data item is assigned a *key* number.



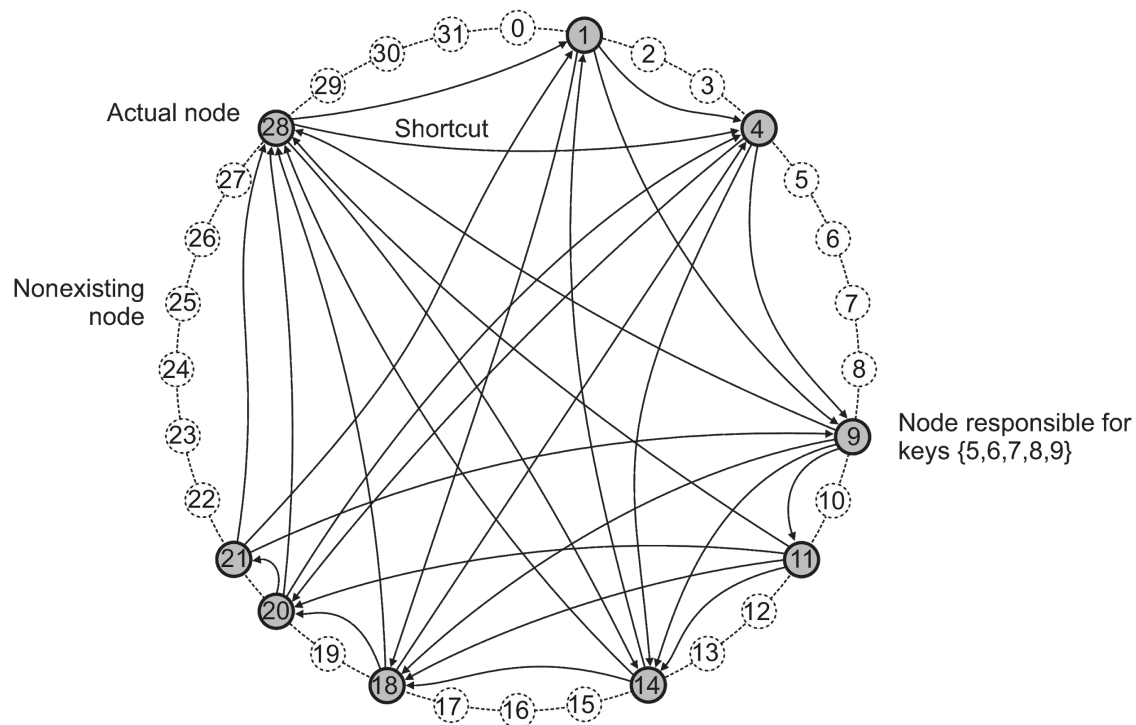
- Data item with key  $k$  is mapped to node with smallest identifier  $p$  greater than or equal to  $k$ .
- This node is referred to as the successor of key  $k$  and denoted  $\text{succ}(k)$ .
- In this example,  $\text{succ}(5) = 9$  and  $\text{succ}(9) = 9$ .




- Each node maintains shortcuts to other nodes.
- This ensures length of the shortest path between any pair of nodes is  $O(\log N)$  where  $N$  is the number of nodes.



- To look up a key, a node will try to forward the request “as far as possible” but without passing it beyond the node responsible for that key.



- Example - node 9 is asked to look up the node responsible for key 3 (which is node 4).





- Node 9 has four shortcuts: to nodes 11, 14, 18, and 28.
- Node 28 is the farthest node and still preceding the one responsible for key 3; it will get the lookup request.
- Node 28 has three shortcuts: to nodes 1, 4, and 14.
- Node 1 is the farthest node and still preceding the one responsible for key 3; it will get the lookup request.
- Node 1 knows that its successor in the ring is node 4, and thus that this is the node responsible for key 3, to which it will subsequently forward the request.


# Searching in unstructured P2P using flooding

- In flooding an issuing node passes a request for a data item to **all neighbours**.
- A receiving node forwards the request to **all** its own neighbours.
- Searching terminates when the data is found or maximum number of hops (time-to-live or TTL) is reached.










# Searching in unstructured P2P using random walk


- In random walk an issuing node passes a request for a data item to a **randomly chosen neighbour**.
- A receiving node forwards the request to a **randomly chosen neighbour**.
- Searching terminates when the data is found or maximum number of hops (time-to-live or TTL) is reached.
- To decrease waiting time, an issuer can start multiple random walks simultaneously.




# Hierarchically organized P2P architecture

- In unstructured peer-to-peer systems, locating data can be slow in large networks (not scalable).
- One solution is to use **super peers** that maintain an index of data items on a subset of nodes.
- Every regular peer, called a **weak peer**, is connected to a super peer.
- All communication from and to a weak peer proceeds through that peer's associated super peer.
- Reduces the number of hops required.


*Super peers*




*Weak peers*



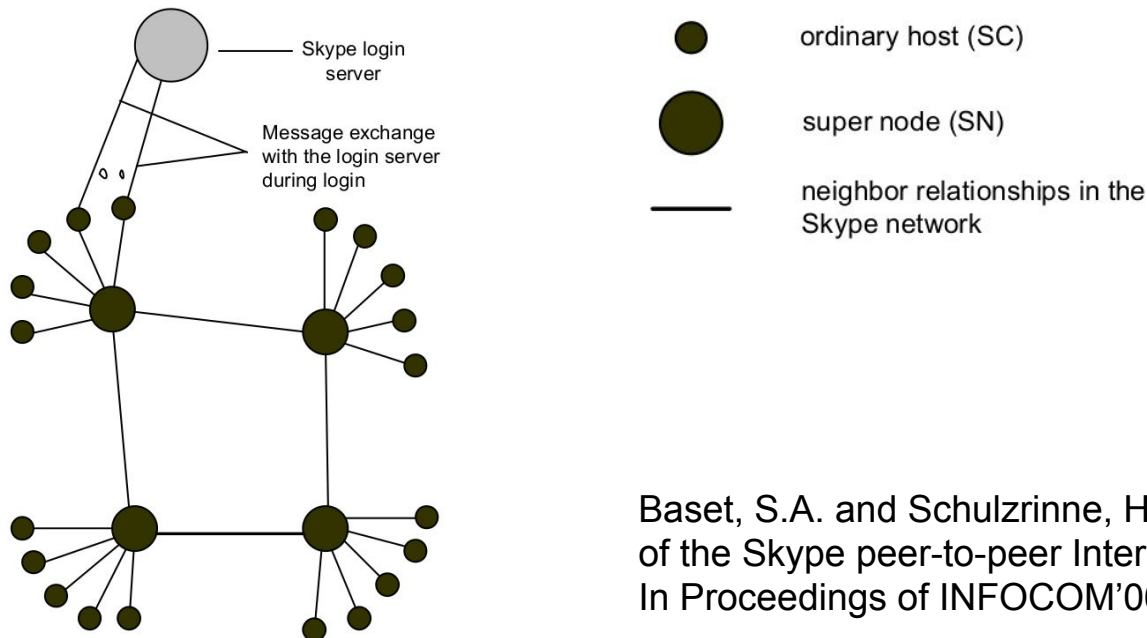






# Skype

- Skype is an example of a hierarchically organized peer-to-peer architecture.
- There is an additional Skype login server.



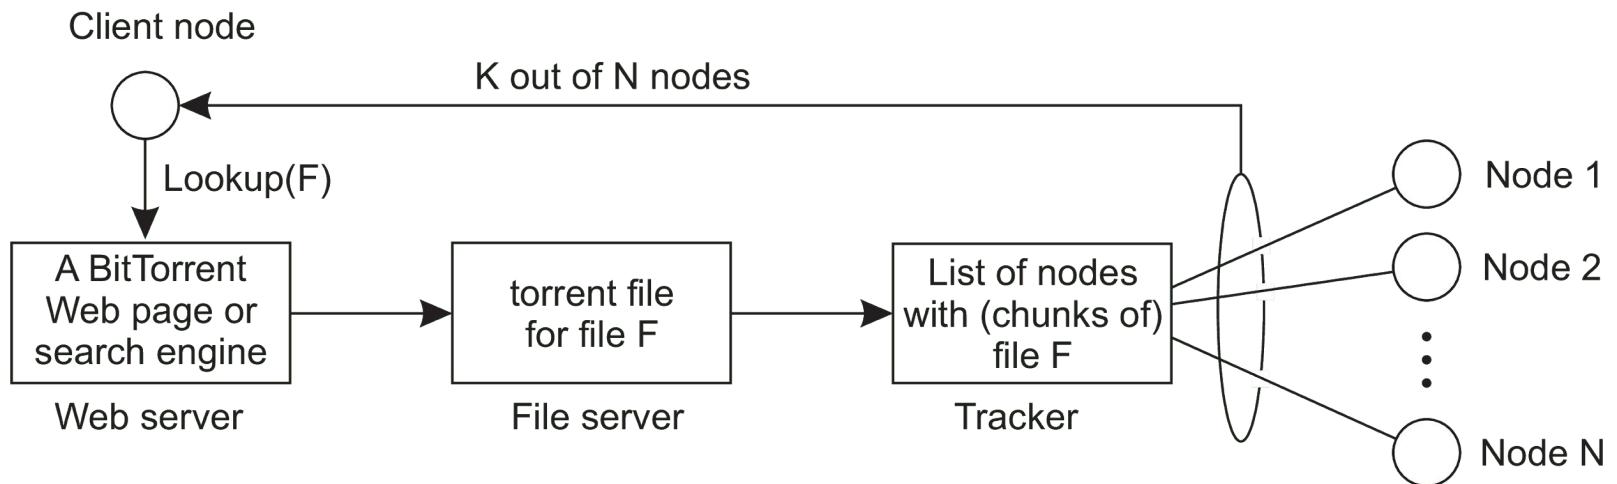
Baset, S.A. and Schulzrinne, H.G. (2006). An analysis of the Skype peer-to-peer Internet telephony protocol. In Proceedings of INFOCOM'06, pp. 1–11.

# Structured Vs Unstructured P2P

	Advantage	Disadvantage
Structured	Guaranteed to locate objects, with time and complexity bounds; relatively low message overhead	Need to build and maintain complex overlay structures
Unstructured	Self-organizing and resilient to node failure	No guarantee on locating objects; prone to excessive message overhead which can affect scalability

# Hybrid architecture - BitTorrent

- User downloads chunks of a file from other users and assemble together yielding the complete file.
- To get started a traditional client-server scheme is deployed.
- A tracker is a server that keeps an accurate account of active nodes that have (chunks of) the requested file.




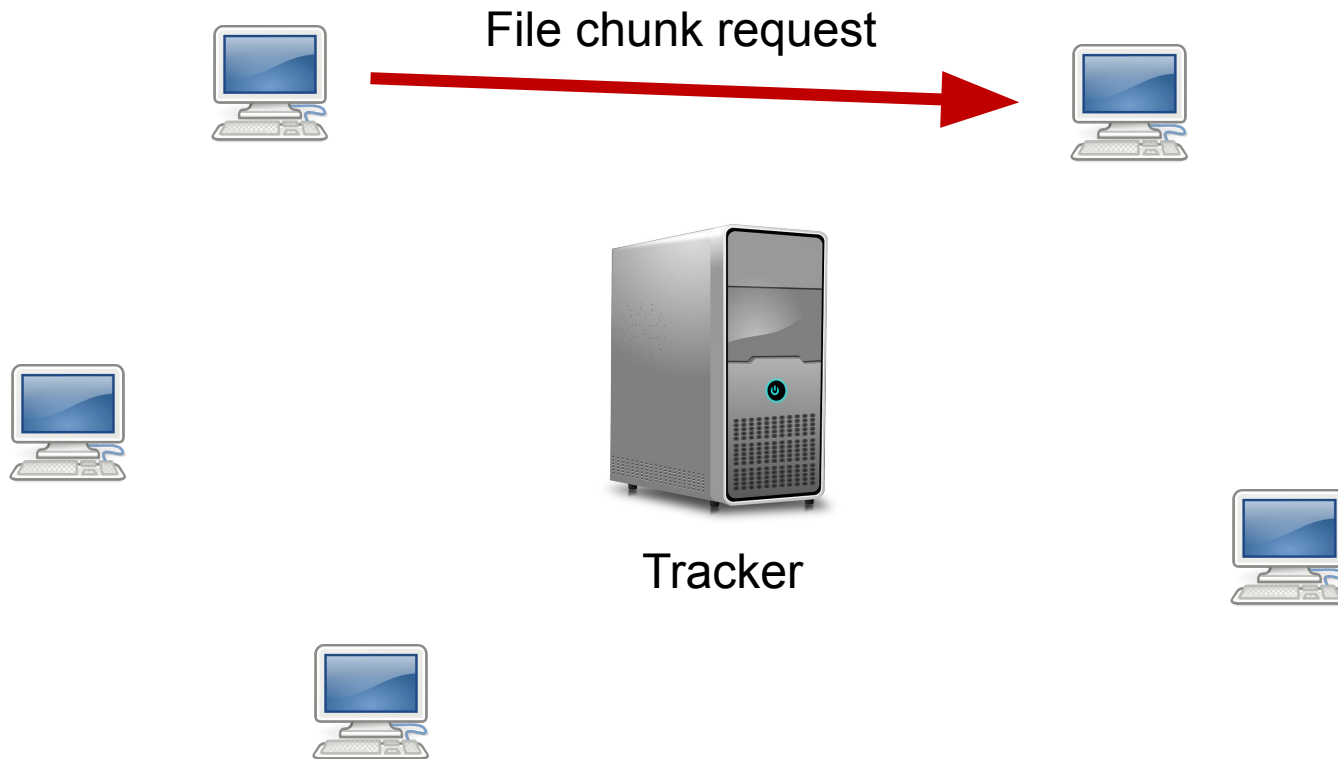
*The principal working of BitTorrent; torrent file contains a link to a tracker.*



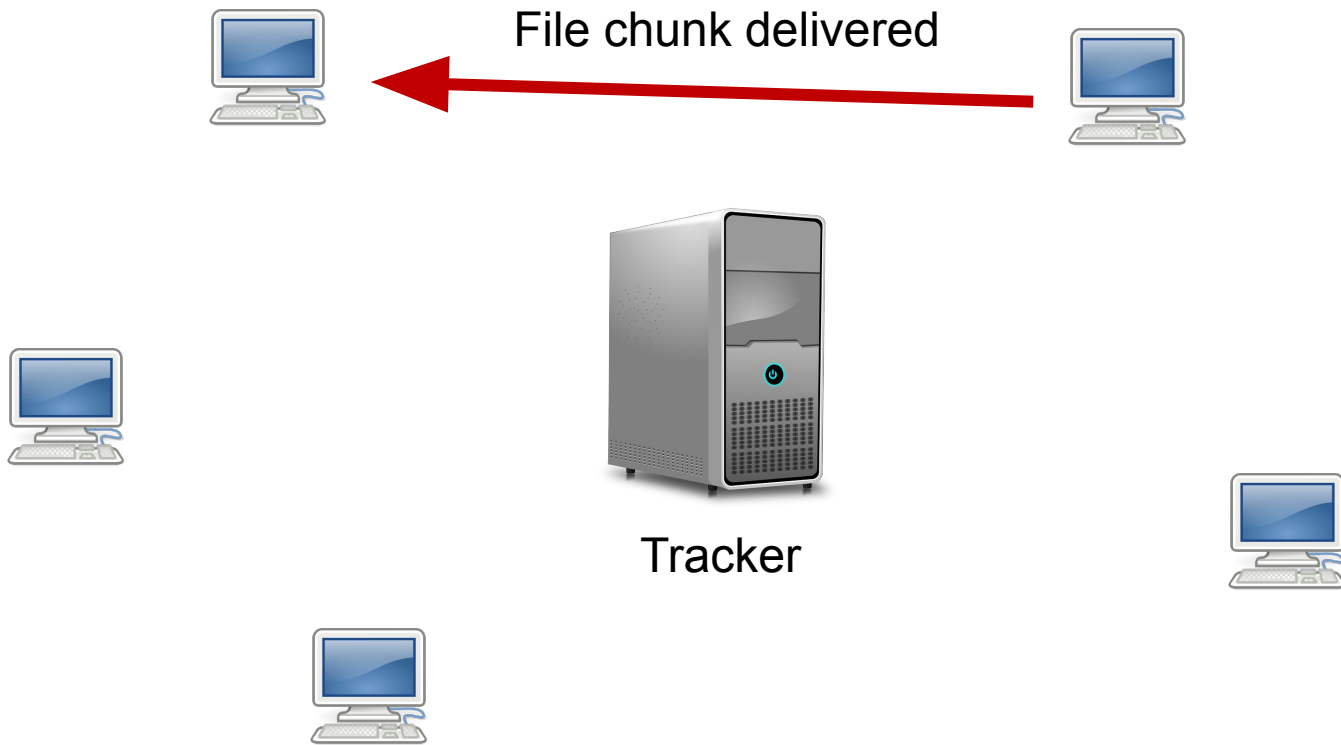
Tracker















# Edge computing

- Computation is performed by devices at the “edge” of the network.
- Reduces the amount of centralized computation.



*Organising a sensor network database, while storing and processing data (a) only at the operator's site or (b) only at the sensors.*