

SECURITY

Dr. Padraig Corcoran

Security in distributed systems consists of two parts:

- **Communication security** - concerns communication between users/processes (principal mechanism - secure channel)
- **Authorisation security** - concerns processes only getting those access rights to resources they are entitled to (principal mechanism - access control).

Introduction to security

- A **security incident** is an event that occurs when an organisation's systems or data is compromised.
- A **security threat** is a potential cause of a security incident.
- A **risk** refers to the potential loss or damage resulting from a security incident.

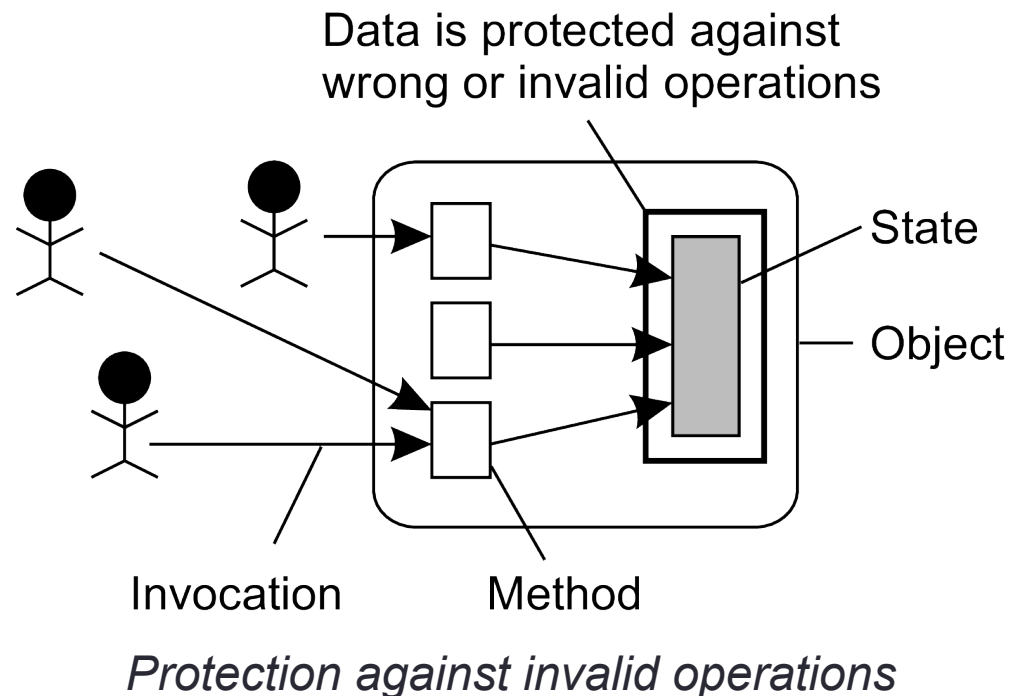
- There are four types of security threats to consider:
 - Interception - unauthorised party gains access to a service or data.
 - Interruption - services or data become unavailable, unusable, destroyed etc.
 - Modification - unauthorised changing of data or tampering with a service so that it no longer adheres to its original specifications.
 - Fabrication - additional data or activity are generated that would normally not exist.

- A **security policy** describes the actions entities in a system are allowed to take and which are prohibited.
- Entities include users, services, data, machines, etc.
- If enforced effectively this will reduce security threats.
- **Security mechanisms** are used to enforce security policies.


- Important security mechanisms are:
 - Encryption - transforms data into something an attacker cannot understand.
 - Authentication - used to verify the claimed identity of a user, client, server, host, or other entity.
 - Authorisation - check whether a client is authorised to perform the action requested.
 - Auditing - used to trace which clients accessed what, and in which way (does not provide protection but is useful).

Focus of control

- Three approaches can be considered when protecting a system from security threats.
- First approach - security policy concentrates directly on the protection of data associated with the application.




- Second approach - security policy concentrates on specifying which operations may be invoked, and by whom, when certain data/resources are to be accessed.



Protection against unauthorized invocations.

- Third approach - security policy concentrates on users by specifying who has access to the application, irrespective of operations they want to carry out.

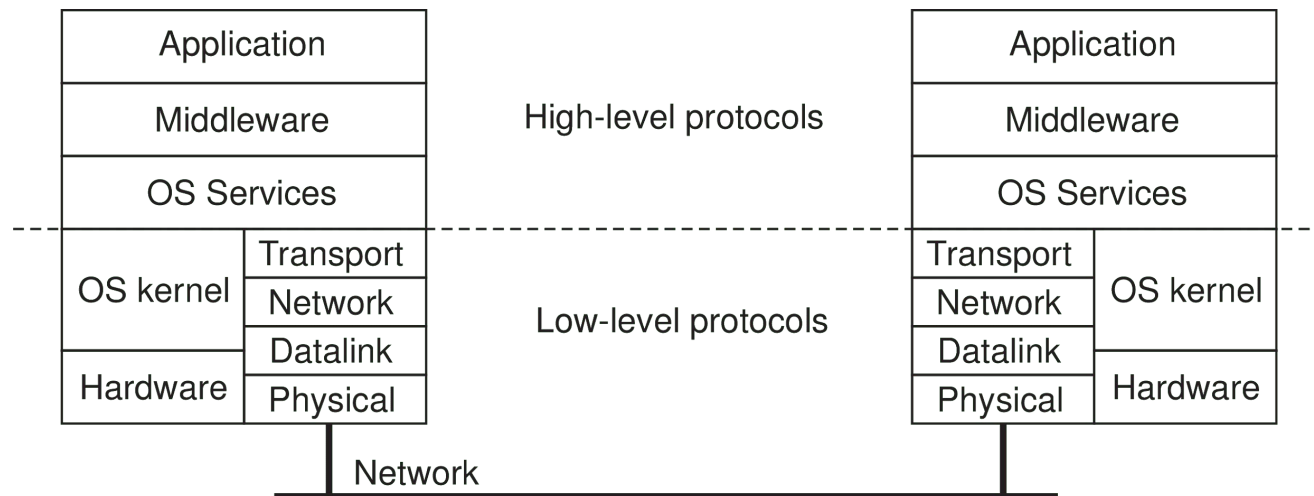


Data is protected by
checking the role of invoker

Protection against unauthorised invocations.

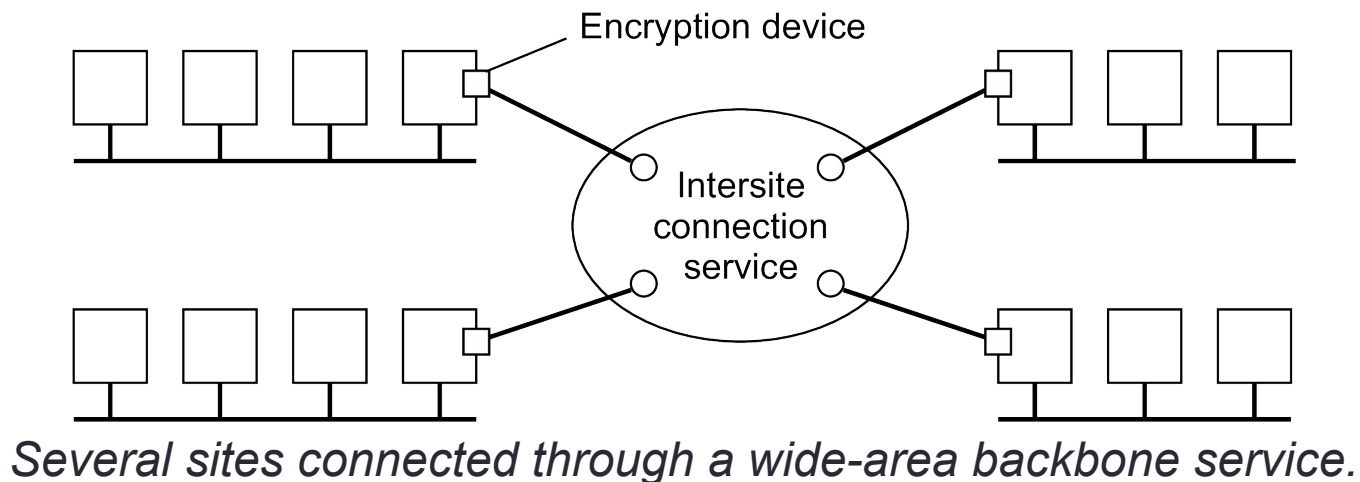
Layering of security mechanisms

- Given organisation of a system into layers, it must be decided at which level to place security mechanisms.
- In which layer security mechanisms are placed depends on trust in how secure the services are in a layer.



The logical organisation of a distributed system into several layers.

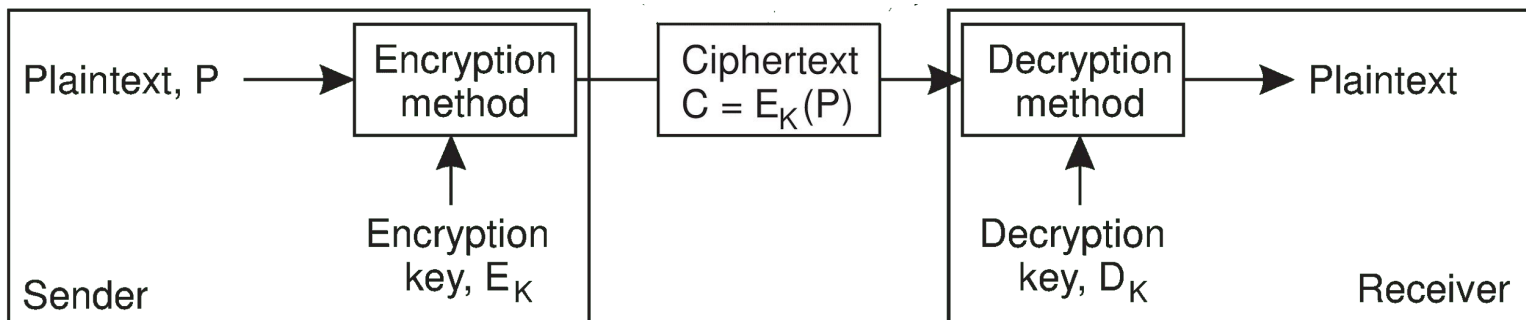
- Consider organisation located at different sites connected through a backbone connecting various LANS.
- Security provided by placing encryption devices at each backbone switch.
- If security of inter-site traffic is not trusted, one may use a **Transport Layer Security (TLS)** service.



Secure communication channel using cryptography

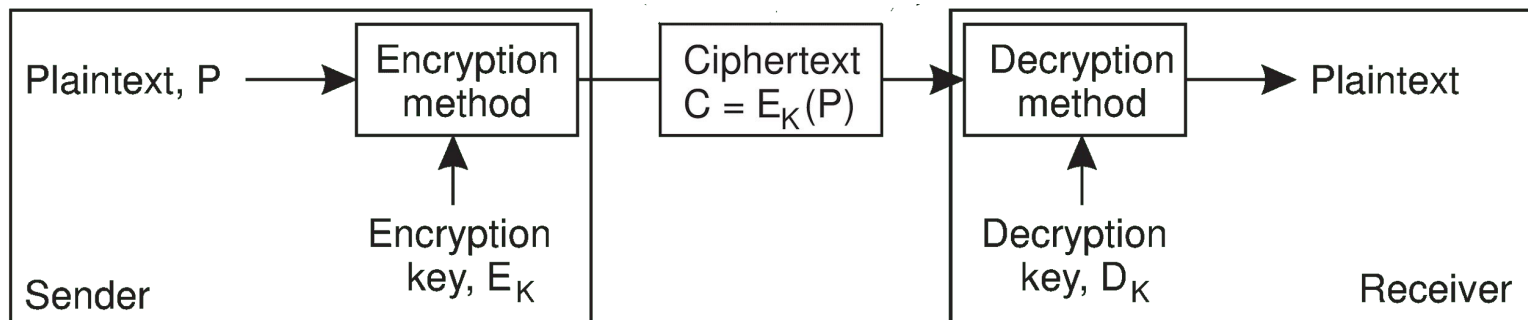
- Consider a sender S wanting to transmit a message m to receiver R .
- To protect against security threats, the sender **encrypts** m into an unintelligible message m' , and subsequently sends m' to R .
- R , in turn, must **decrypt** the received message m' into its original form m .

- Encryption and decryption are accomplished by using cryptographic methods parameterised by keys.
- The original form of the message sent is called the **plaintext** and encrypted form is called the **ciphertext**.



P and C equal the plaintext and ciphertext respectively.

- $C = E_K(P)$ denotes that the ciphertext C is obtained by encrypting the plaintext P using the key K .
- $P = D_K(C)$ denotes decryption of the ciphertext C using key K , resulting in the plaintext P .



P and C equal the plaintext and ciphertext respectively.

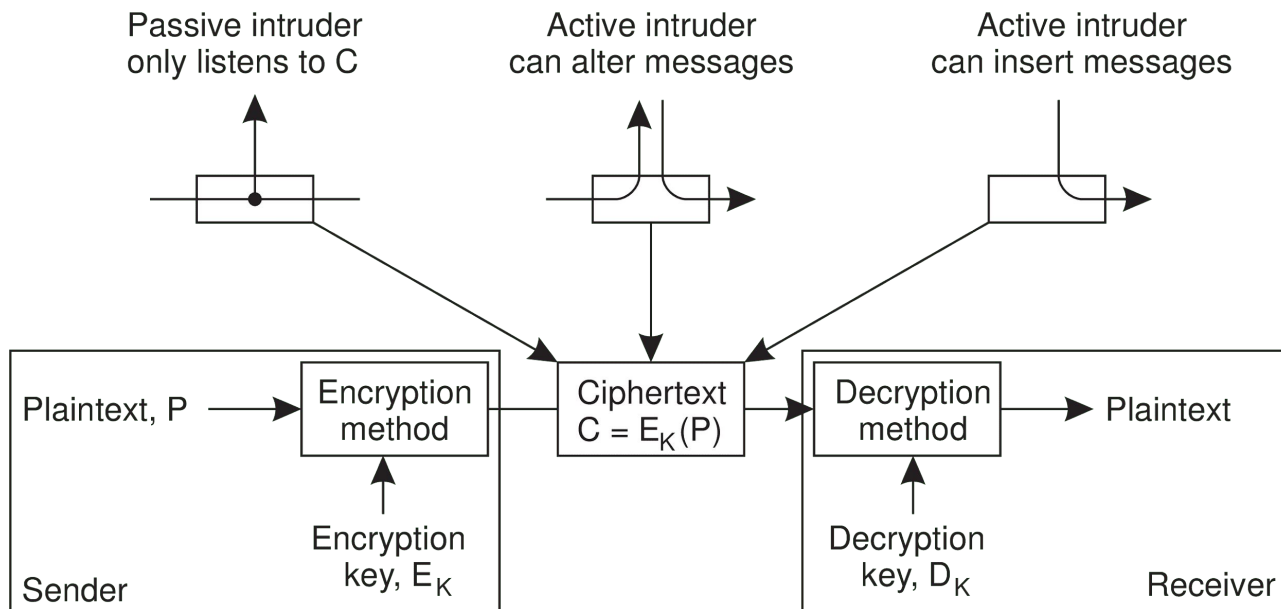
- Caesar's Method is a simple encryption method.
- Each letter of the English alphabet is substituted by a new letter k places ahead of it.
- When $k = 3$, 'A' and 'Y' in plaintext become 'D' and 'B' respectively.
- Formally, we can represent the encryption and decryption mechanisms as:

$$E_k: L \rightarrow (L+k) \bmod 26$$

$$D_k: L \rightarrow (L-k) \bmod 26$$

where L denotes the numerical index of a letter.

- There are three attacks for which encryption helps:
 - Message is intercepted without either sender or receiver being aware that eavesdropping is happening.
 - The message is modified.
 - Message inserted into the communication system, attempting to make receiver believe this message came from sender.



Intruders and eavesdroppers in communication.

- A **cryptosystem** consists of three algorithms:
 - One for key generation.
 - One for encryption.
 - One for decryption.
- Distinction between cryptosystems based on whether or not the encryption and decryption keys are equal.

- In a **symmetric (secret-key) cryptosystem**, the same key is used to encrypt and decrypt a message:

$$P = D_K(E_K(P))$$

- Recall, E_K denotes encryption with key K , D_K denotes decryption with key K and P denotes the plaintext.
- Sender and receiver required to share the same key and this shared key must be kept secret.

- In an **asymmetric (public-key) cryptosystem**, the keys are different but form a unique pair:

$$P = D_{KD}(E_{KE}(P))$$

- Key pair KE and KD are used for encryption and decryption respectively.
- One of the keys in an asymmetric cryptosystem is kept private; the other is made public.
- Which one of the encryption or decryption keys that is actually made public depends on how the keys are used.

Example

- Alice wants to send a confidential message to Bob.
- She uses Bob's public key to encrypt the message.
- Bob is the only one holding the associated and private decryption key.
- Therefore, he is also the only person that can decrypt the message.

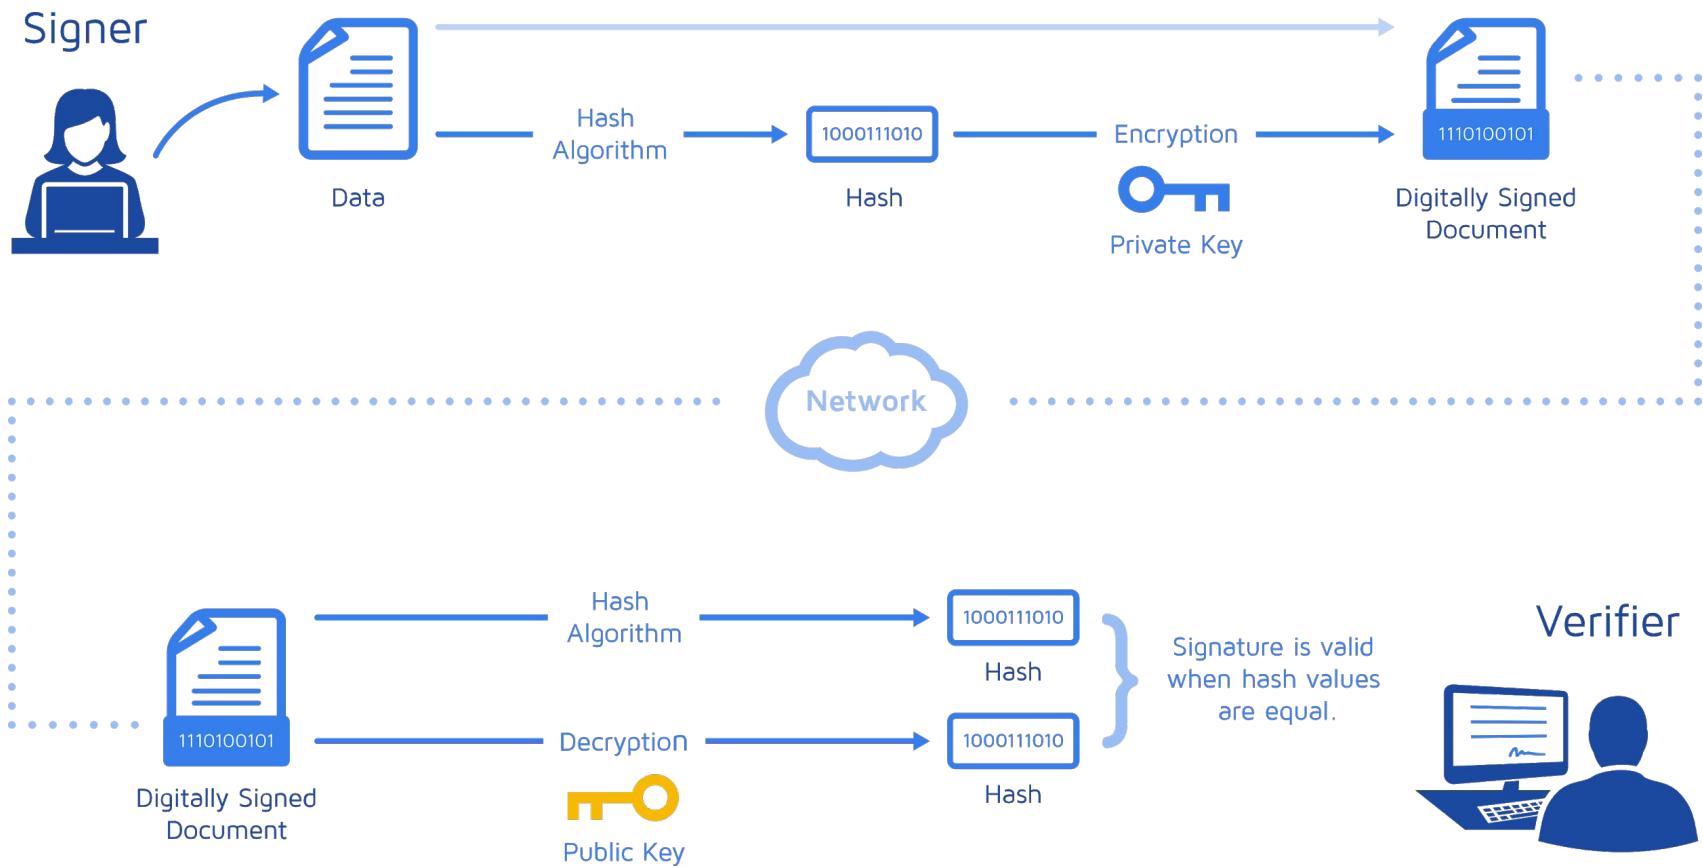
Example

- Bob wants to know for sure that the message he just received actually came from Alice. That is, perform **message authentication**.
- Alice can keep her encryption key private to encrypt the messages she sends.
- If Bob can decrypt a message using Alice's public key the message must have come from Alice, because the decryption key is uniquely tied to the encryption key.

Digital signatures

- Bob agreed to sell an item to Alice for \$500 via e-mail.
- Two issues that need to be considered regarding the integrity of the message are:
 - Alice needs to be assured that Bob will not change the \$500 mentioned in her message into something higher.
 - Bob needs to be assured that Alice cannot deny ever having sent the message, for example, because she had second thoughts.
- Both issues can be dealt with if Alice **digitally signs** the message such that her signature is uniquely tied to its content.
- **Digital signatures** may be implemented using a public key cryptosystems.

- A hash function $H(.)$ reduces a variable length input to a fixed length output.
- A hash function will have the following properties:
 - Computed $m = H(M)$ is easy but computing the inverse is impossible.
 - If $M \neq M'$ then $H(M) \neq H(M')$ (strong collision resistance).



<https://www.docusign.co.uk/how-it-works/electronic-signature/digital-signature/digital-signature-faq>

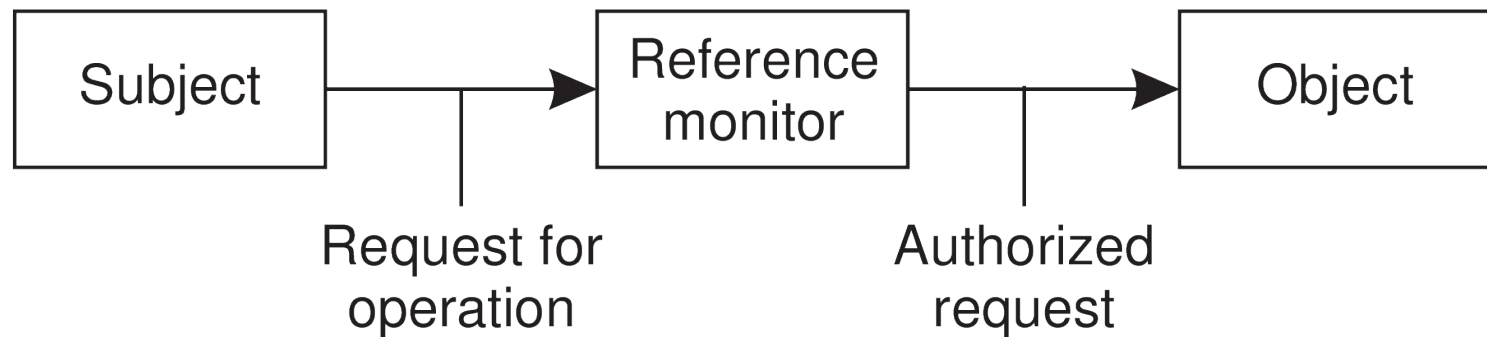
Theorem - If the two hashes are equal the message was sent by the person corresponding to the private-public key pair in question.

Proof - Someone encrypted the hash using their private. The private key in question belongs to that person. Therefore they are the only person who could have sent the message.

Access control

- Consider a client requesting that a remote object method is invoked.
- Such a request should be carried out only if the client has sufficient **access rights** for that invocation.
- Verifying access rights is known as **access control**.
- **Authorisation** concerns granting access rights.

- **Access control** concerns protecting an object against invocations by subjects that are not allowed to have specific (or even any) of the methods carried out.
- Commonly enforced using a **reference monitor** which
 - Records which subject may do what.
 - Decides whether a subject is allowed to have a specific operation carried out.



General model of controlling access to objects.

Access control matrix

- Access rights of subjects with respect to objects may be modelled using an **access control matrix**.
- Each subject is represented by a row in this matrix; each object is represented by a column.
- Given a matrix M , the entry $M[s,o]$ lists operations subject s can request to be carried out on object o .

- Subjects: p, q
- Objects: f, g, h
- Operations: r, w, x, a, o


	f	g	h
p	rwo	r	$rwXO$
q	a		

Access control matrix

- Many entries in the matrix will be empty; a single subject will have access to relatively few objects.

Firewall

- A **firewall** acts as a **reference monitor** between a distributed system and other distributed systems.
- All messages are routed through the reference monitor and inspected before they are passed through.
- Unauthorized traffic is discarded and not allowed to continue.



A common implementation of a firewall.

- Two main types of firewalls - **packet-filtering gateway** and **application-level gateway**.
- Packet-filtering gateway - makes decisions based on the source and destination address in a packet header.
- Example usage - protect an internal server against requests from external clients.
- Application-level gateway - inspects the content of an incoming or outgoing message.
- Example usage - filtering spam e-mail.

Denial of service (DoS) attack

- An attack related to access control is preventing authorized processes from accessing resources.
- Usually implemented by an attacker sending a larger number of fake requests to a server.
- DoS attacks that come from a single process can often be handled effectively.
- In a **distributed denial of service (DDoS)** a set of processes jointly attempt to bring down a service.

- In many cases the attackers hijack a large group of machines which unknowingly participate in the attack.
- Such a set/network of machines is known as a **botnet**.

Secure mobile code

- Mobile code introduces a number of security threats; malicious code may corrupt its host.
- Represents an **access control** problem; program should not be allowed unauthorised access to host's resources.
- Security policies require we can allow access to local resources in a flexible, yet fully controlled manner.

- Security mechanisms to deal with threats and enforce policies:
 1. Execute code in a restricted environment called a **sandbox**.
 2. Ensure code **authenticity** and **integrity**; enforce a **security policy** based on where the program came from. Only trusted code is accepted.
- Authenticity - assures users that they know where the code came from.
- Integrity - verifies that the code hasn't been tampered with since its publication.
- Note, **code signing** may be used to ensure code authenticity and integrity.

Secure naming

- When a client retrieves an object based on a name, how does it know that it got back the correct object?
- We have three issues to worry about:
 - Validity: is the object returned a complete, unaltered copy of what was stored at the server?
 - Provenance: can the server that returned the object be trusted as a genuine supplier?
 - Relevance: is what was returned relevant considering what was asked?

- One solution to secure naming is to securely bind the name of an object to its content through hashing.
- That is, take the hash function $H(O)$ as the name of the object O .
- Given this, one can check that the object has not be altered by comparing its name to its hash.
- This is a form of a **self-certifying name**.